

README

Perry Kundert

November 29, 2014

Contents

1	Comm. Protocol Python Parser and Originator	2
1.1	Installing	2
1.1.1	Installing from source	2
1.1.2	Python Version and OS Support	3
2	Protocols	3
2.1	EtherNet/IP CIP Logix Controller Simulator and Client . . .	3
2.1.1	EtherNet/IP Controller Simulator	4
2.1.2	EtherNet/IP Controller I/O Customization	5
2.1.3	EtherNet/IP Controller Client	6
2.1.4	EtherNet/IP Controller Client I/O Python API	8
2.1.5	Web Interface	9
3	Deterministic Finite Automata	14
3.1	Basic State Machines	15
3.2	Composite Machines	15
3.3	Machines from Regular Expressions	16
3.3.1	Unicode Support	16
4	Running State Machines	17
5	Historical	17
5.1	The timestamp	17
6	Virtualization	18
6.1	Vagrant	18
6.1.1	Building a Vagrant Image	19
6.2	Docker	19

1 Comm. Protocol Python Parser and Originator

Cpppo (pronounced ‘c’+3*’p’+’o’ in Python) is used to implement binary communications protocol parsers. The protocol’s communication elements are described in terms of state machines which change state in response to input events, collecting the data and producing output data artifacts.

1.1 Installing

Cpppo depends on several Python packages:

Package	For?	Description
greenery	all	Regular Expression parsing and state machinery library
web	web API	The web.py HTTP web application framework
pytz	history	The Python time-zone library
tzlocal	history	Access to system’s local timezone (on Mac, Windows)
pymodbus	remote	Modbus/TCP support for polling Schneider compatible PLCs
pytest	unit test	A Python unit-test framework

To install ‘cpppo’ and its required dependency ‘greenery’ using pip:

```
$ pip install cpppo
```

To install all required and optional Python modules, use:

```
pip install -r requirements.txt
pip install -r requirements-optional.txt
```

If you need system admin privileges to install packages, you may need to use ‘sudo’ on the above pip commands:

```
$ sudo pip install ...
```

1.1.1 Installing from source

Clone the repo by going to your preferred source directory and using:

```
$ git clone git@github.com:pjkundert/cpppo.git
```

You can then install from the provided setuptools-based setup.py installer:

```
$ cd cpppo
$ python setup.py install
```

1.1.2 Python Version and OS Support

Cpppo is implemented and fully tested on both Python 2 (2.6 and 2.7), and Python 3 (3.3 and 3.4). However, some of cpppo’s modules are not (yet) fully supported in both versions:

- The pymodbus module does not support Python 3, so Modbus/TCP support for polling remote PLCs is only available for Python 2.
- Greenery supports both Python 2 and 3, but doesn’t provide meaningful Unicode (UTF-8) support in Python 2, so regular expression based DFAs are only supported for Python 3.

Linux (native or Docker containerized), Mac and Windows OSs are supported. However, Linux or Mac are recommended for stability, performance and ease of use.

2 Protocols

The protocols implemented are described here.

2.1 EtherNet/IP CIP Logix Controller Simulator and Client

A subset of the EtherNet/IP client and server protocol is implemented, and a simulation of a subset of a Allen-Bradley ControlLogix 5561 Controller is provided. It is capable of simulating ControlLogix Tag access, via the Read/Write Tag [Fragmented] services.

Only EtherNet/IP “Unconnected” type connections are supported. These are (somewhat anomalously) a persistent connection to a single EtherNet/IP device (such as a Controller), which allow a sequence of CIP service requests (commands) to be sent to arbitrary CIP objects resident on the target device.

A Tag is simply a shortcut to a specific EtherNet/IP CIP Object Instance and Attribute. Instead of the Client needing to know the specific Instance and Attribute numbers, the more easily remembered and meaningful Tag may be supplied in the request path.

2.1.1 EtherNet/IP Controller Simulator

To run a simulation of a subset of a ControlLogix(tm) Controller with the tag 'SCADA' for you to read/write, run:

```
$ python -m cpppo.server.enip --print SCADA=INT[1000]
```

Alternatively, invoke the supplied wrapper script:

```
enip_server --print SCADA=INT[1000]
```

This is especially handy under Python 2.6, where you must manually specify the module entry point, eg. `cpppo.server.enip.__main__`.

The following options are available when you execute the `cpppo.server.enip` module:

Specify a different local interface and/or port to bind to (default is `:44818`, indicating all interfaces and port 44818):

```
-a|--address [<interface>][:<port>]
```

Change the verbosity (supply more to increase further):

```
-v[vv...]|--verbose
```

Specify a constant or variable delay to apply to every response, in fractional seconds:

```
-d|--delay #.#[-#.#]
```

Specify an HTTP web server interface and/or port, if a web API is desired (just `.'` will enable the web API on defaults `:80`, or whatever interface was specified for `-address`):

```
-w|--web [<interface>][:<port>]
```

To send log output to a file (limited to 10MB, rotates through 5 copies):

```
-l|--log <file>
```

To print a summary of PLC I/O to stdout:

```
-p|--print
```

You may specify as many tags as you like on the command line; at least one is required:

```
<tag>=<type>[<length>] # eg. SCADA=INT[1000]
```

The available types are INT (16-bit), SINT (8-bit) DINT (32-bit) integer, and REAL (32-bit float).

2.1.2 EtherNet/IP Controller I/O Customization

If you require access to the read and write I/O events streaming from client(s) to and from the EtherNet/IP CIP Attributes hosted in your simulated controller, you can easily make a custom `cpppo.server.enip.device` Attribute implementation which will receive all PLC Read/Write Tag [Fragmented] request data.

We provide two examples; one which records a history of all read/write events to each Tag, and one which connects each Tag to the current temperature of the city with the same name as the Tag.

- Record Tag History

For example purposes, we have implemented the `cpppo.server.enip.historize` module which intercepts all I/O (and exceptions) and writes it to the file specified in the first command-line argument to the module:

```
$ python -m cpppo.server.enip.historize some_file.hst Tag_Name=INT[1000] &
$ tail -f some_file.txt
# 2014-07-15 22:03:35.945: Started recording Tag: Tag_Name
2014-07-15 22:03:44.186 ["Tag_Name", [0, 3]]      {"write": [0, 1, 2, 3]}
...
```

(in another terminal)

```
$ python -m cpppo.server.enip.client Tag_Name[0-3]=[0,1,2,3]
```

You can examine the code in `cpppo/server/enip/historize.py` to see how to easily implement your own customization of the EtherNet/IP CIP Controller simulator.

If you invoke the ‘main’ method provided by `cpppo.server.enip.main` directly, all command-line args will be parsed, and the EtherNet/IP service will not return control until termination. Alternatively, you may start the service in a separate `threading.Thread` and provide it with a list of configuration options. Note that each individual EtherNet/IP Client session is serviced by a separate Thread, and thus all method invocations arriving at your customized Attribute object need to process data in a Thread-safe fashion.

- City Temperature Tag

In this example, we intercept read requests to the Tag, and look up the current temperature of the city named with the Tag’s name. This example is simple enough to include here (see `cpppo/server/enip/weather.py`):

```

import sys, logging, json
try:
    from urllib2 import urlopen
except ImportError:
    from urllib.request import urlopen

from cpppo.server.enip import device
from cpppo.server.enip.main import main

class Attribute_weather( device.Attribute ):
    def __getitem__( self, key ):
        try:
            url = "http://api.openweathermap.org/data/2.5/weather?units=metric&q="
            weather = json.loads( urlopen( url ).read() )
            return [ weather['main']['temp'] ]
        except Exception as exc:
            logging.warning( "Couldn't get temperature for %s: %s", self.name, exc )
            raise

    def __setitem__( self, key, value ):
        raise Exception( "Changing the weather isn't that easy..." )

sys.exit( main( attribute_class=Attribute_weather ) )

```

By providing a specialized implementation of `device.Attribute`'s `__getitem__` (which is invoked each time an `Attribute` is accessed), we arrange to query the city's weather at the given URL, and return the current temperature. Of course, `__setitem__` (which would be invoked whenever someone wishes to change the city's temperature) would have a much more complex implementation, the details of which are left as an exercise to the reader...

2.1.3 EtherNet/IP Controller Client

A simple EtherNet/IP Controller Client is provided. Presently, it only knows how to Register, and issue Read Tag Fragmented requests (and only for offset 0; the first block of results):

```
python -m cpppo.server.enip.client SCADA[1]=99 SCADA[0-10]
```

Alternatively, invoke the supplied wrapper script:

```
enip_client SCADA[1]=99 SCADA[0-10]
```

Specify a different local interface and/or port to connect to (default is :44818):

```
-a|--address [<interface>][:<port>]
```

On Windows systems, you must specify an actual interface. For example, if you started the cpppo.server.enip simulator above (running on the all interfaces by default), use `--address localhost`.

Change the verbosity (supply more to increase further):

```
-v[vv...]|--verbose
```

Change the default response timeout

```
-t|--timeout #
```

Specify a number of times to repeat the specified operations:

```
-r|--repeat #
```

To send log output to a file (limited to 10MB, rotates through 5 copies):

```
-l|--log <file>
```

To print a summary of PLC I/O to stdout:

```
-p|--print
```

To force use of the Multiple Service Packet request, which carries multiple Read/Write Tag [Fragmented] requests in a single EtherNet/IP CIP I/O operation (default is to issue each request as a separate I/O operation):

```
-m|--multiple
```

To force the client to use plain Read/Write Tag commands (instead of the Fragmented commands, which are the default):

```
-n|--no-fragment
```

You may specify as many tags as you like on the command line; at least one is required. An optional register (range) can be specified (default is register 0):

```
<tag> <tag>[<reg>] <tag>[<reg>-<reg>] # eg. SCADA SCADA[1] SCADA[1-10]
```

Writing is supported; the number of values must exactly match the data specified register range:

```
<tag>=<value>                                # scalar, eg. SCADA=1
<tag>[<reg>-<reg>]=<value>,<value>,...      # vector range
<tag>[<reg>]=<value>                          # single element of a vector
<tag>[<reg>-<reg>]=(DINT)<value>,<value>      # cast to SINT, INT, DINT or REAL
```

If any <value> contains a ‘.’ (eg. ‘9.9,10’), all values are deemed to be REAL; otherwise, they are integers and default to a type INT. To force a specific type (and limit the values to the appropriate value range), you may specify a “cast” to a specific type, eg. ‘TAG[4-6]=(INT)1,2,3’. The types SINT, INT, DINT and REAL are supported.

In addition to symbolic Tag addressing, numeric Class/Instance/Attribute addressing is available. A Class, Instance and Attribute address values are in decimal by default, but hexadecimal, octal etc. are available using escapes, eg. 26 == 0x1A == 0o49 == 0b100110.

```
@<class>/<instance>/<attribute>              # read a scalar, eg. @0xFF/01/0x1A
@<class>/<instance>/<attribute>[99]=1        # write element, eg. @511/01/26=1
```

2.1.4 EtherNet/IP Controller Client I/O Python API

Dispatching a multitude of EtherNet/IP CIP I/O operations to a Controller is very simple (see `cpppo/server/enip/client/io.py`):

```
host                = 'localhost'    # Controller IP address
port                = address[1]      # default is port 44818
depth               = 1               # Allow 1 transaction in-flight
multiple            = 0               # Don't use Multiple Service Packet
fragment            = False           # Don't force Read/Write Tag Fragmented
timeout             = 1.0             # Any PLC I/O fails if it takes > 1s
printing            = True            # Print a summary of I/O
tags                = ["Tag[0-9]+16=(DINT)4,5,6,7,8,9", "@0x2/1/1", "Tag[3-5]"]

connection          = client.connector( host=host, port=port, timeout=timeout )
operations           = client.parse_operations( tags )
failures,transactions = connection.process(
    operations=operations, depth=depth, multiple=multiple,
    fragment=fragment, printing=printing, timeout=timeout )

sys.exit( 1 if failures else 0 )
```


Try it out by starting up a simulated Controller:

```
$ python -m cpppo.server.enip Tag=DINT[10] &  
$ python -m cpppo.server.enip.io
```

The API is able to “pipeline” requests – issue multiple requests on the wire, while simultaneously harvesting prior requests. This is absolutely necessary, in order to obtain reasonable I/O performance over high-latency links (eg. via Satellite).

To use it, create a `client.connector(host=..., port=...)` which establishes and registers a CIP connection to a Controller. Then, produce a sequence of operations (eg, parsed from “Tag[0-9]+16=(DINT)5,6,7,8,9” or from numeric Class, Instance and Attribute numbers “@2/1/1”), and dispatch the requests using connector methods `.pipeline` or `.synchronous` (to access the details of the requests and the harvested replies), or `.process` to simply get a summary of I/O failures and total transactions.

2.1.5 Web Interface

The following actions are available via the web interface. It is designed to be primarily a REST-ful HTTP API returning JSON, but any of these requests may be made via a web browser, and a minimal HTML response will be issued.

Start a Logix Controller simulator on port 44818 (the default), with a web API on port 12345:

```
python -m cpppo.server.enip -v --web :12345 SCADA=INT[1000]
```

The api is simple: `api/<group>/<match>/<command>/<value> .` There are 3 groups: “options”, “tags” and “connections”. If you don’t specify `<group>` or `<match>`, they default to the wildard “*”, which matches anything.

So, to get everything, you should now be able to hit the root of the api with a browser at: `http://localhost:12345/api`, or with `wget` or `curl`:

```
$ wget -qO - http://localhost:12345/api  
$ curl http://localhost:12345/api
```

and you should get something like:

```
$ curl http://localhost:12345/api
```

```

{
  "alarm": [],
  "command": {},
  "data": {
    "options": {
      "delay": {
        "value": 0.0
      }
    },
    "server": {
      "control": {
        "disable": false,
        "done": false,
        "latency": 5.0,
        "timeout": 5.0
      }
    },
    "tags": {
      "SCADA": {
        "attribute": "SCADA          INT[1000] == [0, 0, 0, 0, 0, 0,...]",
        "error": 0
      }
    }
  },
  "since": null,
  "until": 1371731588.230987
}

```

- options/delay/value

To access or modify some specific thing in the matching object(s), add a <command> and <value>:

```

$ curl http://localhost:12345/api/options/delay/value/0.5
{
  "alarm": [],
  "command": {
    "message": "options.delay.value=u'0.5' (0.5)",
    "success": true
  },
  "data": {

```

```

        "options": {
            "delay": {
                "value": 0.5
            }
        },
        "since": null,
        "until": 1371732496.23366
    }

```

It will perform the action of assigning the `<value>` to all of the matching `<command>` entities. In this case, since you specified a precise `<group>` “options”, and `<match>` “delay”, exactly one entity was affected: “value” was assigned “0.5”. If you are running a test client against the simulator, you will see the change in response time.

As a convenience, you can use `/<value>` or `=<value>` as the last term in the URL:

```

$ curl http://localhost:12345/api/options/delay/value/0.5
$ curl http://localhost:12345/api/options/delay/value=0.5

```

- `api/options/delay/range`

If you’ve started the simulator with `-delay=0.1-0.9` (a delay range), you can adjust this range to a new range, using:

```

$ curl http://localhost:12345/api/options/delay/range=0.5-1.5

```

You can cause it to never respond (in time), to cause future connection attempts to fail:

```

$ curl http://localhost:12345/api/options/delay/value=10.0

```

Or, if you’ve configured a delay range using `-delay=#-#`, use:

```

$ curl http://localhost:12345/api/options/delay/range=10.0-10.0

```

Restore connection responses by restoring a reasonable response timeout.

- `api/server/control/done` or `disable`

To prevent any future connections, you can (temporarily) disable the server, which will close its port (and all connections) and await further instructions:

```
$ curl http://localhost:12345/api/server/control/disable/true
```

Re-enable it using:

```
$ curl http://localhost:12345/api/server/control/disable/false
```

To cause the server to exit completely (and of course, causing it to not respond to future requests):

```
$ curl http://localhost:12345/api/server/control/done/true
```

- `api/server/control/latency` or `timeout`

The default socket I/O blocking ‘latency’ is .1s; this is the time it may take for each existing connection to detect changes made via the web API, eg. signalling EOF via `api/connections/eof/true`. The ‘timeout’ on each thread responding defaults to twice the latency, to give the thread’s socket I/O machinery time to respond and then complete. These may be changed, if necessary, if simulation of high-latency links (eg. satellite) is implemented (using other network latency manipulation software).

- `api/tags/<tagname>/error`

To force all successful accesses to a certain tag (eg. SCADA) to return a certain error code, you can set it using:

```
$ curl http://localhost:12345/api/tags/SCADA/error=8
```

Restore it to return success:

```
$ curl http://localhost:12345/api/tags/SCADA/error/0
```

- `api/tags/<tagname>/attribute[x]`

To access or change a certain element of a tag, access its attribute at a certain index (curl has problems with this kind of URL):

```
wget -q0 - http://localhost:12345/api/tags/SCADA/attribute[3]=4
```

You can access any specific value to confirm:

```
wget -q0 - http://localhost:12345/api/tags/SCADA/attribute[3]
{
  "alarm": [],
  "command": {
    "message": "tags.SCADA.attribute[2]: 0",
    "success": true
  },
  "data": {
    "tags": {
      "SCADA": {
        "attribute": "SCADA          INT[1000] == [0, 0, 0, 4, 0, 0,
        ...]",
        "error": 0
      }
    }
  },
  "since": null,
  "until": 1371734234.553135
}
```

- `api/connections/*/eof`

To immediately terminate all connections, you can signal them that they've experienced an EOF:

```
$ curl http://localhost:12345/api/connections/*/eof/true
```

If there are any matching connections, all will be terminated. If you know the port and IP address of the interface from which your client is connecting to the simulator, you can access its connection specifically:

```
$ curl http://localhost:12345/api/connections/10_0_111_121_60592/eof/true
```

To wait for all connections to close, you can issue a request to get all connections, and wait for the 'data' attribute to become empty:

```
$ curl http://localhost:12345/api/connections
{
```

```

    "alarm": [],
    "command": {},
    "data": {
        "connections": {
            "127_0_0_1_52590": {
                "eof": false,
                "interface": "127.0.0.1",
                "port": 52590,
                "received": 1610,
                "requests": 17
            },
            "127_0_0_1_52591": {
                "eof": false,
                "interface": "127.0.0.1",
                "port": 52591,
                "received": 290,
                "requests": 5
            }
        }
    },
    "since": null,
    "until": 1372889099.908609
}
$ # ... wait a while (a few tenths of a second should be OK)...
$ curl http://localhost:12345/api/connections
{
    "alarm": [],
    "command": null,
    "data": {},
    "since": null,
    "until": 1372889133.079849
}


```

3 Deterministic Finite Automata

A `cpppo.dfa` will consume symbols from its source iterable, and yield (machine, state) transitions ‘til a terminal state is reached. If ‘greedy’, it will transition ‘til we reach a terminal state and the next symbol does not produce a transition.

For example, if 'abbb,ab' is presented to the following machine with a no-input state E, and input processing states A and (terminal) B, it will accept 'ab' and terminate, unless greedy is specified in which case it will accept 'abbb' and terminate.

3.1 Basic State Machines



This machine is easily created like this:

```


# Basic DFA that accepts ab+
E = cpppo.state( "E" )
A = cpppo.state_input( "A" )
B = cpppo.state_input( "B", terminal=True )
E['a'] = A
A['b'] = B
B['b'] = B

BASIC = cpppo.dfa( 'ab+', initial=E, context='basic' )

```

3.2 Composite Machines

A higher-level DFA can be produced by wrapping this one in a cpppo.dfa, and giving it some of its own transitions. For example, lets make a machine that accepts 'ab+' separated by '[]*'.
'[]*'



This is implemented:


```
# Composite state machine accepting ab+, ignoring ,[ ]* separators
ABP          = cpppo.dfa( "ab+", initial=E, terminal=True )
SEP          = cpppo.state_drop( "SEP" )
ABP[' ','']  = SEP
SEP[' ' '']  = SEP
SEP[None]    = ABP

CSV          = cpppo.dfa( 'CSV', initial=ABP, context='csv' )
```

When the lower level state machine doesn't recognize the input symbol for a transition, the higher level machine is given a chance to recognize them; in this case, a `' '` followed by any number of spaces leads to a `state_drop` instance, which throws away the symbol. Finally, it uses an “epsilon” (no-input) transition (indicated by a transition on `None`) to re-enter the main CSV machine to process subsequent symbols.

3.3 Machines from Regular Expressions

We use <https://github.com/ferno/greenery> to convert regular expressions into `greenery.fsm` machines, and post-process these to produce a `cpppo.dfa`. The regular expression `'(ab+)((,[]*)(ab+))*'` is equivalent to the above (except that it doesn't ignore the separators), and produces the following state machine:



A regular expression based `cpppo.dfa` is created thus:

```
# A regular expression; the default dfa name is the regular expression itself.
REGEX          = cpppo.regex( initial='(ab+)((,[ ]*)(ab+))*', context='regex' )
```

3.3.1 Unicode Support

Cppo supports Unicode (UTF-8) on both Python 2 and 3. However, `greenery` provides meaningful Unicode support only under Python 3. Therefore,

if you wish to use Unicode in regular expressions, you must use Python 3.

4 Running State Machines

State machines define the grammar for a language which can be run against a sentence of input. All these machines ultimately use `state_input` instances to store their data; the path used is the `cpppo.dfa`'s `<context> + '_input'`:

```
data = cpppo.dotdict()
for machine in [ BASIC, CSV, REGEX ]:
    path = machine.context() + '.input' # default for state_input data
    source = cpppo.peekable( str( 'abbbb, ab' ) )
    with machine:
        for i,(m,s) in enumerate( machine.run( source=source, data=data ) ):
            print( "%s #%3d; next byte %3d: %-10.10r: %r" % (
                m.name_centered(), i, source.sent, source.peek(), data.get(path) ) )
        print( "Accepted: %r; remaining: %r\n" % ( data.get(path), ''.join( source ) ) )
print( "Final: %r" % ( data ) )
```

5 Historical

Recording and playing back time series data is often required for industrial control development and testing. Common pain points are:

- time stamp formats, especially if timezone information is required
- storage/access of time series data, which may be compressed
- playback of the data at various speeds

The `cpppo.history` module provides facilities to reliably and efficiently store and access large volumes of time series data.

5.1 The timestamp

Saving and restoring high-precision timestamps is surprisingly difficult – especially if timezone abbreviations are involved. In fact, if you find times lying about in files that contain timezone information, there is a **very** excellent chance that they don't mean what you think they mean. However, it is universally necessary to deal in dates and times in a user's local timezone; it

is simply not generally acceptable to state times in UTC, and expect users to translate them to local times in their heads.

The `cpppo.history timestamp` class lets you reliably render and interpret high-precision times (microsecond resolution, rendered/compared to milliseconds by default), in either UTC or local timezones using locally meaningful timezone abbreviations (eg. ‘MST’ or ‘MDT’), instead of the globally unambiguous but un-intuitive full timezone names (eg. ‘Canada/Mountain’ or ‘America/Edmonton’).

6 Virtualization

Software with an interface acting as a PLC is often deployed as an independent piece of infrastructure with its own IP address, etc. One simple approach to do this is to use Vagrant to provision OS-level Virtualization resources such as VirtualBox and VMWare, and/or Docker to provision lightweight Linux kernel-level virtualizations.

Using a combination of these two facilities, you can provision potentially hundreds of “independent” PLC simulations on a single host – each with its own IP address and configuration.

6.1 Vagrant

If you are not running on a host capable of directly hosting Docker images, one can be provided for you. Install Vagrant (<http://vagrantup.com>) on your system, and then use the `cpppo/GNUMakefile` target to bring up a VirtualBox or VMWare Fusion (license required: <http://www.vagrantup.com/vmware>):

```
$ make vmware-debian-up # or virtualbox-ubuntu-up
```

Connect to the running virtual machine:

```
$ make vmware-debian-ssh
...
vagrant@jessie64:~$
```

Both Debian and Ubuntu Vagrantfiles are provided, which produce a VM image capable of hosting Docker images. Not every version is available on every platform, depending on what version of VMware or Virtualbox you are running; see the `GNUMakefile` for details.

6.1.1 Building a Vagrant Image

The Debian Jessie + Docker VirtualBox and VMware images used by the Vagrantfiles are hosted at <http://box.hardconsulting.com>. When you use the cpppo/GNUMakefile targets to bring up a Vagrant box (eg. ‘make virtualbox-debian-up’), the appropriate box is downloaded using ‘vagrant box add ...’. If you don’t trust these boxes (the safest position), you can rebuild them yourself.

- Packer.io

Using the packer tool, build a VirtualBox (or VMware) image. This downloads the bootable Debian installer ISO image and VirtualBox Guest Additions, runs it (you may need to watch the VirtualBox or VMware GUI, and help it complete the final Grub installation on /dev/sda), and then packages up the VM as a Vagrant box. We’ll rename it jessie64, and augment the zerodisk.sh script to flush its changes to the device:

```
$ cd src/cpppo/packer
$ make vmware-jessie64 # or virtualbox-jessie64
...
```

Once it builds successfully, add the new box to the ../docker/debian Vagrant installation, to make it accessible:

```
$ make add-vmware-jessie64 # or add-virtualbox-jessie64
```

Now, you can fire up the new VirtualBox image using Vagrant, and the targets provided in the cpppo/GNUMakefile:

```
$ cd src/cpppo
$ make vmware-debian-up
```

6.2 Docker

We’ll assume that you now have a prompt on a Docker-capable machine. Start a Docker container using the pre-built cpppo/cpppo image hosted at <https://index.docker.io/u/cpppo/>. This will run the image, binding port 44818 on localhost thru to port 44818 on the running Docker image, and will run the cpppo.server.enip module with 1000 16-bit ints on Tag “SCADA”:

```
$ docker run -p 44818:44818 -d cpppo/cpppo python -m cpppo.server.enip SCADA=dint[
6da5183740b4
$
```

A canned Docker image is provided which automatically runs an instance of `cpppo.server.enip` hosting the “SCADA=dint¹” tag by default (you can provide alternative tags on the command line, if you wish):

```
$ docker run -p 44818:44818 -d cpppo/scada
```

Assuming you have `cpppo` installed on your local host, you can now test this. We’ll read a single value and a range of values from the tag `SCADA`, repeating 10 times:

```
$ python -m cpppo.server.enip.client -r 10 SCADA[1] SCADA[0-10]
10-08 09:40:29.327 ... SCADA[ 1-1 ] == [0]
10-08 09:40:29.357 ... SCADA[ 0-10 ] == [0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0]
10-08 09:40:29.378 ... SCADA[ 1-1 ] == [0]
10-08 09:40:29.406 ... SCADA[ 0-10 ] == [0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0]
10-08 09:40:29.426 ... SCADA[ 1-1 ] == [0]
10-08 09:40:29.454 ... SCADA[ 0-10 ] == [0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0]
10-08 09:40:29.476 ... SCADA[ 1-1 ] == [0]
10-08 09:40:29.503 ... SCADA[ 0-10 ] == [0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0]
10-08 09:40:29.523 ... SCADA[ 1-1 ] == [0]
10-08 09:40:29.551 ... SCADA[ 0-10 ] == [0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0]
10-08 09:40:29.571 ... SCADA[ 1-1 ] == [0]
10-08 09:40:29.600 ... SCADA[ 0-10 ] == [0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0]
10-08 09:40:29.622 ... SCADA[ 1-1 ] == [0]
10-08 09:40:29.648 ... SCADA[ 0-10 ] == [0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0]
10-08 09:40:29.669 ... SCADA[ 1-1 ] == [0]
10-08 09:40:29.697 ... SCADA[ 0-10 ] == [0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0]
10-08 09:40:29.717 ... SCADA[ 1-1 ] == [0]
10-08 09:40:29.745 ... SCADA[ 0-10 ] == [0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0]
10-08 09:40:29.769 ... SCADA[ 1-1 ] == [0]
10-08 09:40:29.796 ... SCADA[ 0-10 ] == [0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0]
10-08 09:40:29.796 ... Client ReadFrg. Average 20.266 TPS ( 0.049s ea).
$
```

¹DEFINITION NOT FOUND: 1000

6.2.1 Creating Docker images from a Dockerfile

Get started by going to `.../cpppo/docker/cpppo/cpppo/Dockerfile` on your local machine. If you started a Vagrant VM from this directory (eg. `make vmware-up`), this is also mounted inside that machine `/src/cpppo`. Once there, have a look at `docker/cpppo/cpppo/Dockerfile`. If you go into that directory, you can re-create the Docker image:

```
$ cd /src/cpppo/docker/cpppo/cpppo
$ docker build -t cpppo/cpppo .
```

Or, lets use it as a base image for a new Dockerfile. Lets just formalize the command we ran previously so we don't have to remember to type it in. Create a new Dockerfile in, say, `cpppo/docker/cpppo/scada/`:

```
FROM          cpppo/cpppo
MAINTAINER    Whoever You Are "whoever@example.com"
EXPOSE        44818
# We'll always run this as our base command
ENTRYPOINT    [ "python", "-m", "cpppo.server.enip" ]
# But we will allow this to be (optionally) overridden
CMD           [ "SCADA=dint[1000]" ]
```

Then, we can build and save the container under a new name:

```
docker build -t cpppo/scada .
docker run -p 44818
```

This is (roughly) what is implemented in `docker/cpppo/scada/Dockerfile`.